

```
In [17]: import pandas as pd
import numpy as np
import statsmodels.api as sm
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.graphics.regressionplots import influence_plot

# Load the dataset (adjust path if necessary)
data = pd.read_csv('Auto_Data.csv')

# Check the first few rows to ensure the data is loaded correctly
print(data.head())
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130	3504	12.0	70	
1	15.0	8	350.0	165	3693	11.5	70	
2	18.0	8	318.0	150	3436	11.0	70	
3	16.0	8	304.0	150	3433	12.0	70	
4	17.0	8	302.0	140	3449	10.5	70	

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino

```
In [18]: # Assuming 'data' is your DataFrame
# Convert 'horsepower' and 'mpg' to numeric (in case there are non-numeric values)
data['horsepower'] = pd.to_numeric(data['horsepower'], errors='coerce')
data['mpg'] = pd.to_numeric(data['mpg'], errors='coerce')

# Drop rows with missing values
data = data.dropna(subset=['horsepower', 'mpg'])

# Define the predictor (horsepower) and response (mpg)
X = data['horsepower']
y = data['mpg']

# Add a constant (intercept) to the model
X = sm.add_constant(X)

# Fit the model
model = sm.OLS(y, X).fit()

# Print the summary of the regression
print(model.summary())
```

OLS Regression Results

```
=====
=====
Dep. Variable: mpg R-squared: 0.606
Model: OLS Adj. R-squared: 0.605
Method: Least Squares F-statistic: 99.7
Date: Wed, 26 Nov 2025 Prob (F-statistic): 7.03
e-81
Time: 12:45:20 Log-Likelihood: -11
78.7
No. Observations: 392 AIC: 2
361.
Df Residuals: 390 BIC: 2
369.
Df Model: 1
Covariance Type: nonrobust
=====
```

```
=====
=====
          coef    std err      t    P>|t|    [0.025
0.975]
-----
const      39.9359     0.717   55.660    0.000   38.525    4
1.347
horsepower -0.1578     0.006  -24.489    0.000   -0.171
-0.145
=====
=====
Omnibus: 16.432 Durbin-Watson: 0.920
Prob(Omnibus): 0.000 Jarque-Bera (JB): 1
7.305
Skew: 0.492 Prob(JB): 0.00
0175
Kurtosis: 3.299 Cond. No.
322.
=====
=====
```

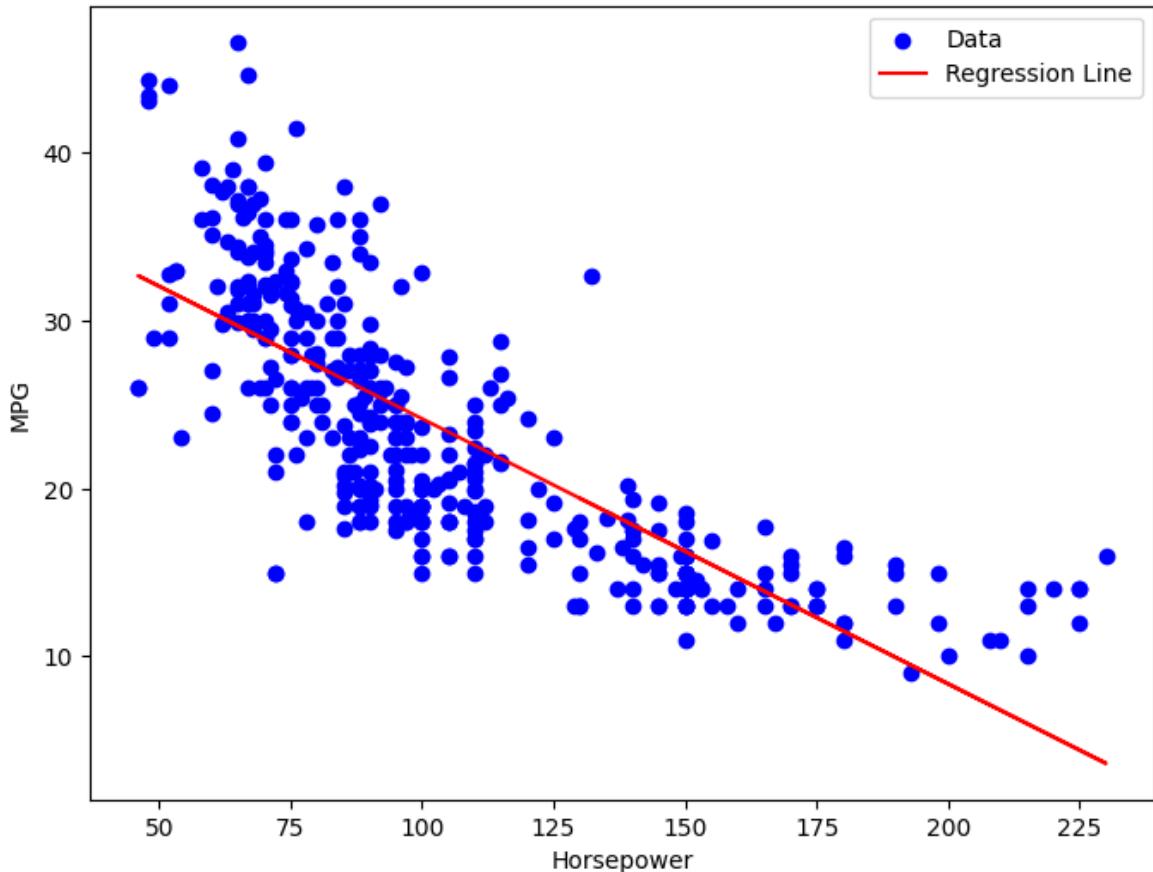
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [19]: # Scatter plot of mpg vs horsepower
fig, ax = plt.subplots(figsize=(8, 6))
ax.scatter(X['horsepower'], y, label="Data", color='blue')

# Plot the regression line (predicted values)
predicted_values = model.predict(X)
ax.plot(X['horsepower'], predicted_values, color='red', label="Regression

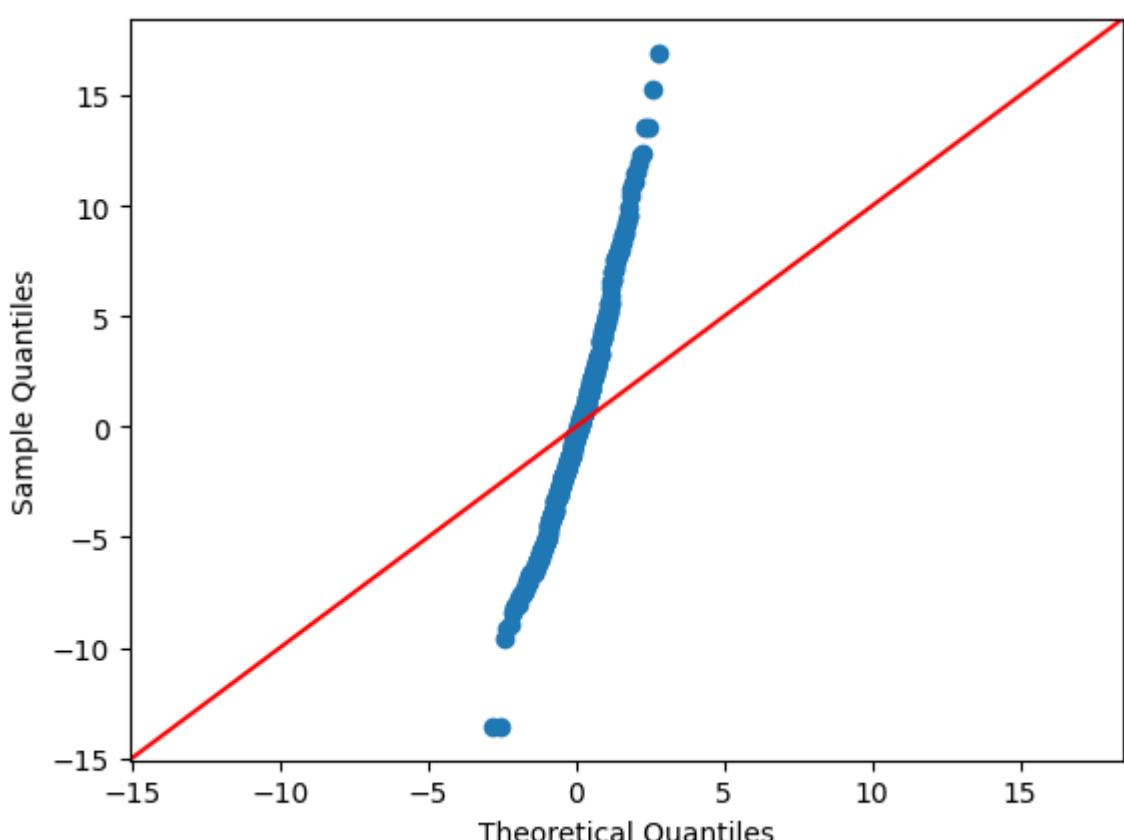
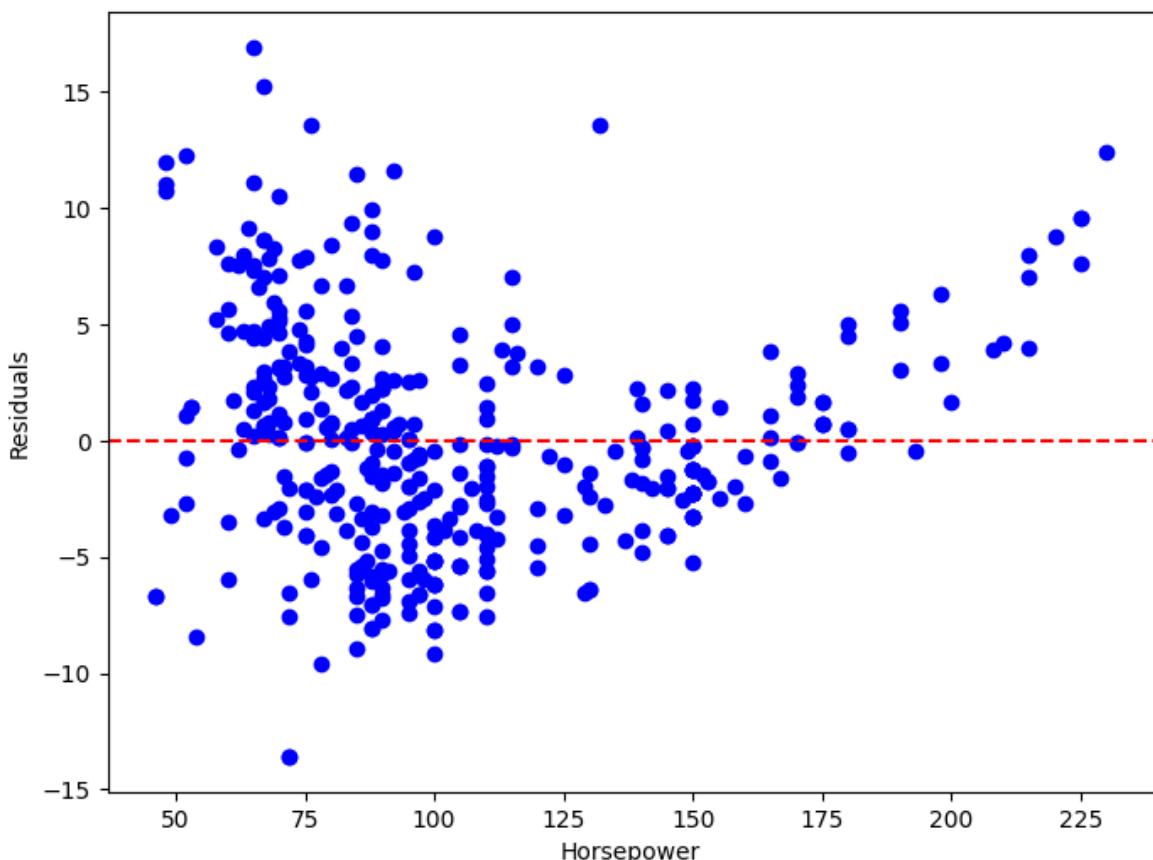
ax.set_xlabel('Horsepower')
ax.set_ylabel('MPG')
ax.legend()
plt.show()
```

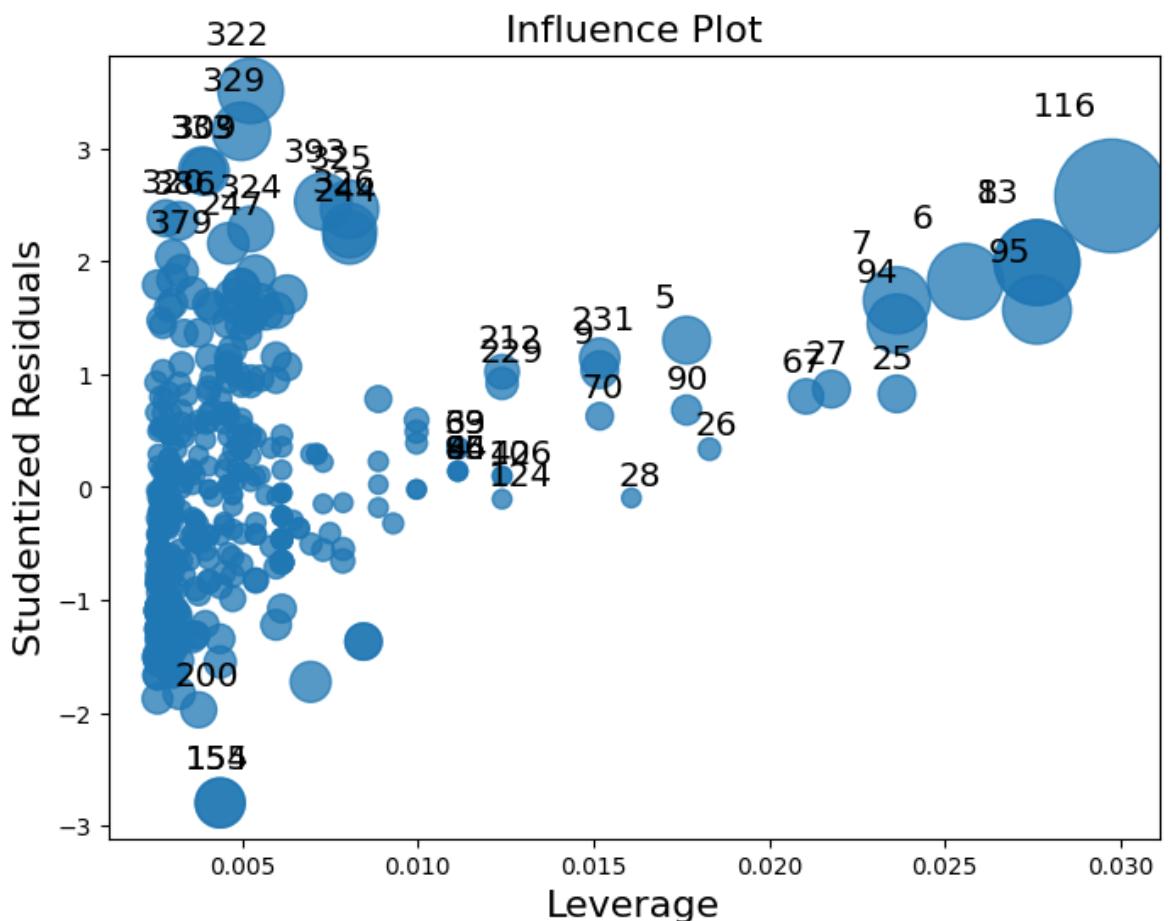


```
In [20]: # Residual plot
residuals = model.resid
fig, ax = plt.subplots(figsize=(8, 6))
ax.scatter(X['horsepower'], residuals, color='blue')
ax.axhline(0, color='red', linestyle='--')
ax.set_xlabel('Horsepower')
ax.set_ylabel('Residuals')
plt.show()

# QQ Plot for Normality
sm.qqplot(residuals, line='45')
plt.show()

# Leverage Plot (to identify influential points)
fig, ax = plt.subplots(figsize=(8, 6))
influence_plot(model, ax=ax)
plt.show()
```





```
In [21]: # Scatterplot matrix using seaborn  
sns.pairplot(data)  
plt.show()
```



```
In [22]: # Select only numeric columns
numeric_data = data.select_dtypes(include=[np.number])

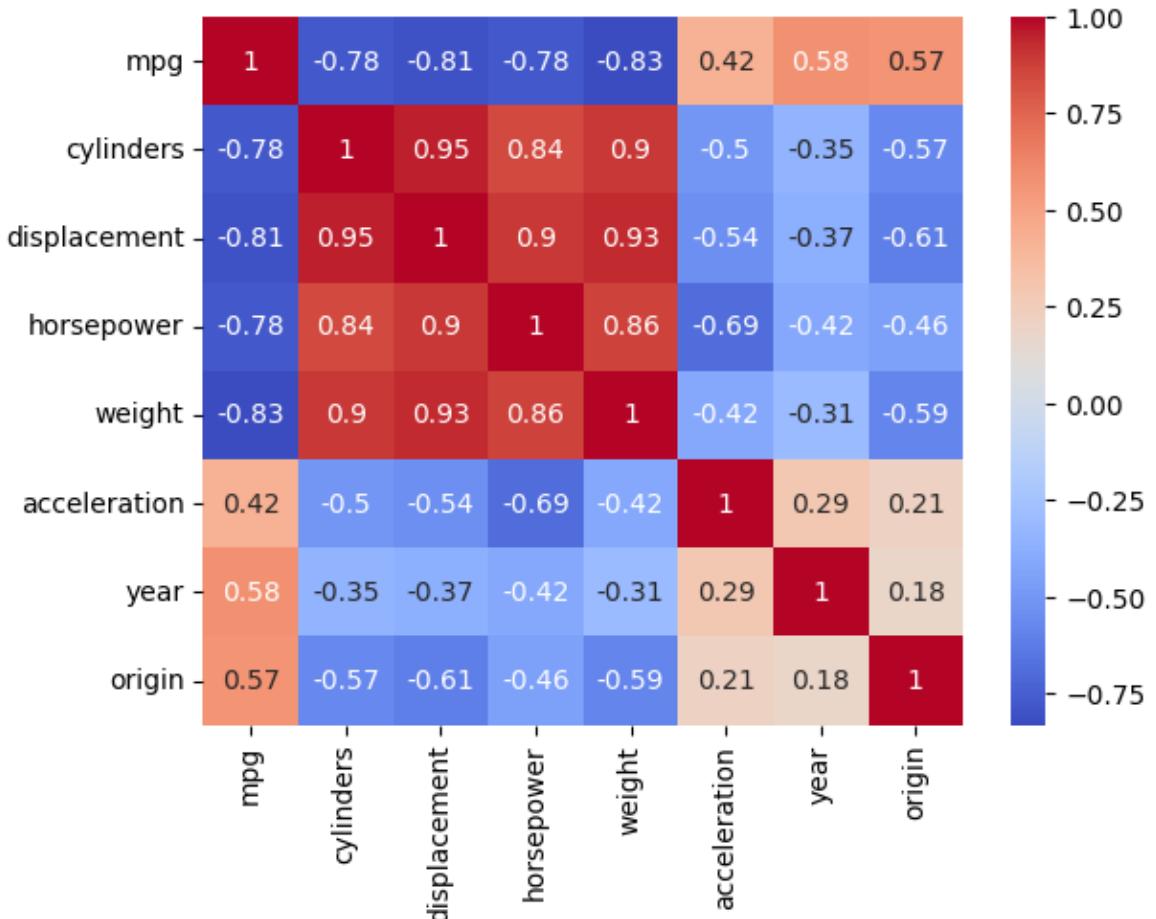
# Calculate the correlation matrix
corr_matrix = numeric_data.corr()

# Print the correlation matrix
print(corr_matrix)

# Plot the correlation matrix using heatmap
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.show()
```

	mpg	cylinders	displacement	horsepower	weight	\
mpg	1.000000	-0.777618	-0.805127	-0.778427	-0.832244	
cylinders	-0.777618	1.000000	0.950823	0.842983	0.897527	
displacement	-0.805127	0.950823	1.000000	0.897257	0.932994	
horsepower	-0.778427	0.842983	0.897257	1.000000	0.864538	
weight	-0.832244	0.897527	0.932994	0.864538	1.000000	
acceleration	0.423329	-0.504683	-0.543800	-0.689196	-0.416839	
year	0.580541	-0.345647	-0.369855	-0.416361	-0.309120	
origin	0.565209	-0.568932	-0.614535	-0.455171	-0.585005	

	acceleration	year	origin
mpg	0.423329	0.580541	0.565209
cylinders	-0.504683	-0.345647	-0.568932
displacement	-0.543800	-0.369855	-0.614535
horsepower	-0.689196	-0.416361	-0.455171
weight	-0.416839	-0.309120	-0.585005
acceleration	1.000000	0.290316	0.212746
year	0.290316	1.000000	0.181528
origin	0.212746	0.181528	1.000000



```
In [23]: # Drop 'mpg' (response) and 'name' (irrelevant feature) from predictors
X_multi = data.drop(columns=['mpg', 'name']) # Exclude 'mpg' and 'name'

# Convert predictors to numeric, ensuring no non-numeric values remain
X_multi = X_multi.apply(pd.to_numeric, errors='coerce')

# Convert response variable 'mpg' to numeric
y_multi = pd.to_numeric(data['mpg'], errors='coerce')

# Handle missing values (if any)
X_multi = X_multi.dropna() # Drop rows with NaN values in predictors
y_multi = y_multi[X_multi.index] # Ensure 'y' aligns with the rows of 'X'
```

```
# Add constant (intercept) to the model  
X_multi = sm.add_constant(X_multi)  
  
# Fit the multiple regression model  
model_multi = sm.OLS(y_multi, X_multi).fit()  
  
# Print the summary of the regression  
print(model_multi.summary())
```

OLS Regression Results

```
=====
=====
Dep. Variable: mpg R-squared: 0.821
Model: OLS Adj. R-squared: 0.818
Method: Least Squares F-statistic: 52.4
Date: Wed, 26 Nov 2025 Prob (F-statistic): 2.04
e-139
Time: 12:45:28 Log-Likelihood: -10
23.5
No. Observations: 392 AIC: 2
063.
Df Residuals: 384 BIC: 2
095.
Df Model: 7
Covariance Type: nonrobust
=====
```

	coef	std err	t	P> t	[0.025
0.975]					
-----	-----	-----	-----	-----	-----
const	-17.2184	4.644	-3.707	0.000	-26.350
-8.087					
cylinders	-0.4934	0.323	-1.526	0.128	-1.129
0.142					
displacement	0.0199	0.008	2.647	0.008	0.005
0.035					
horsepower	-0.0170	0.014	-1.230	0.220	-0.044
0.010					
weight	-0.0065	0.001	-9.929	0.000	-0.008
-0.005					
acceleration	0.0806	0.099	0.815	0.415	-0.114
0.275					
year	0.7508	0.051	14.729	0.000	0.651
0.851					
origin	1.4261	0.278	5.127	0.000	0.879
1.973					

=====

		Durbin-Watson:
0.309	31.906	
Prob(Omnibus):	0.000	Jarque-Bera (JB): 5
3.100		
Skew:	0.529	Prob(JB): 2.95
e-12		
Kurtosis:	4.460	Cond. No. 8.59
e+04		

=====

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 8.59e+04. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [24]: # Example: Add interaction term (horsepower * displacement)
data['horsepower_displacement'] = data['horsepower'] * data['displacement']

# Define predictors including the interaction term
X_interaction = data[['horsepower', 'displacement', 'horsepower_displacement']]
X_interaction = sm.add_constant(X_interaction)

# Fit the model with interaction term
model_interaction = sm.OLS(y_multi, X_interaction).fit()

# Print the summary of the model with interaction term
print(model_interaction.summary())
```

OLS Regression Results

```
=====
=====
Dep. Variable: mpg R-squared: 0.747
Model: OLS Adj. R-squared: 0.745
Method: Least Squares F-statistic: 81.0
Date: Wed, 26 Nov 2025 Prob (F-statistic): 3.00
e-115
Time: 12:45:28 Log-Likelihood: -10
92.1
No. Observations: 392 AIC: 2
192.
Df Residuals: 388 BIC: 2
208.
Df Model: 3
Covariance Type: nonrobust
=====
=====
```

	coef	std err	t	P> t
[0.025 0.975]				
const	53.0511	1.526	34.765	0.000
0.051 56.051				5
horsepower	-0.2343	0.020	-11.960	0.000
-0.273 -0.196				
displacement	-0.0980	0.007	-14.674	0.000
-0.111 -0.085				
horsepower_displacement	0.0006	5.19e-05	11.222	0.000
0.000 0.001				

```
=====
=====
Omnibus: 46.481 Durbin-Watson: 1.057
Prob(Omnibus): 0.000 Jarque-Bera (JB): 7.417
Skew: 0.685 Prob(JB): 1.04
e-19
Kurtosis: 4.865 Cond. No.: 2.47
e+05
=====
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.47e+05. This might indicate that there are strong multicollinearity or other numerical problems.

In [25]: `import numpy as np`

```
# Apply log transformation on horsepower and displacement (since 'weight'
data['log_horsepower'] = np.log(data['horsepower'])
data['log_displacement'] = np.log(data['displacement'])

# Define predictors for the transformed model
```

```
X_transformed = data[['log_horsepower', 'log_displacement']]
X_transformed = sm.add_constant(X_transformed)

# Fit the transformed model
model_transformed = sm.OLS(y_multi, X_transformed).fit()

# Print the summary of the transformed model
print(model_transformed.summary())
```

OLS Regression Results

=====

Dep. Variable: mpg R-squared: 0.724

Model: OLS Adj. R-squared: 0.723

Method: Least Squares F-statistic: 5.10.3

Date: Wed, 26 Nov 2025 Prob (F-statistic): 1.77

e-109

Time: 12:45:28 Log-Likelihood: -11.08.9

No. Observations: 392 AIC: 2.224.

Df Residuals: 389 BIC: 2.236.

Df Model: 2

Covariance Type: nonrobust

=====

	coef	std err	t	P> t	[0.025
0.975]					
-----	-----	-----	-----	-----	-----
const	101.0700	2.915	34.667	0.000	95.338
106.802					
log_horsepower	-9.0199	1.237	-7.289	0.000	-11.453
-6.587					
log_displacement	-7.0675	0.798	-8.860	0.000	-8.636
-5.499					

=====

Omnibus: 30.256 Durbin-Watson: 0.994

Prob(Omnibus): 0.000 Jarque-Bera (JB): 4.2.413

Skew: 0.569 Prob(JB): 6.17

e-10

Kurtosis: 4.140 Cond. No. 103.

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In []: